

Flow matching and Diffusion Models

Lecture 1: Generative AI with SDEs

Section 1: From Generation to Sampling

Data distribution: how "likely" are we to find this picture in the Internet.

translate subjective statement into a statement of probability theory

↑
how good an image is how likely it is under the data distribution

- Data distribution: distribution of objects that we want to generate: p_{data}

- Probability density: $p_{\text{data}}: \mathbb{R}^d \rightarrow \mathbb{R}_{>0}$, $t \mapsto p_{\text{data}}(t)$

Generation means sampling the data distribution

unconditional generation: a model with fixed prompt: eg. prompt: "dog". you always get dog from this model

conditional generation involves sampling from $z \sim p_{\text{data}}(\cdot | y)$, where y is a conditioning variable.
conditional data distribution

our goal is to train a generative model to transform samples from a simple distribution p_{init} (e.g. Gaussian) into samples from p_{data}



look a little different

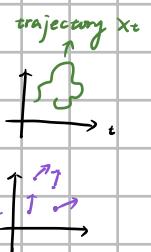
Section 2: Flow and Diffusion Models

Flow Models

fundamental object of flow: ① trajectory → a function of time $X: [0, 1] \rightarrow \mathbb{R}^d$ $t \mapsto \text{vector } X_t$

② vector field → $u: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$ $(x, t) \mapsto u_t(x)$

it gives you directions at every point



ODE: $\dot{x}_t = u(x_t)$ (initial condition)

$\frac{dx}{dt} = u(x_t) \Rightarrow$ the direction of x_t is given by the location of x_t , and then specified by u .

flow: a collection of trajectories that follow the ODE, collection of solutions to an ODE for a lot of initial condition. we gather a lot of solutions for different initial conditions, and then we gather them all in one function and call that a flow

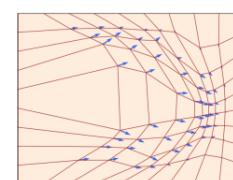
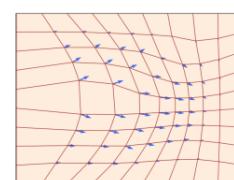
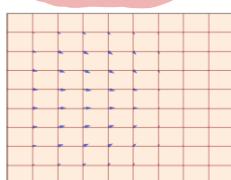
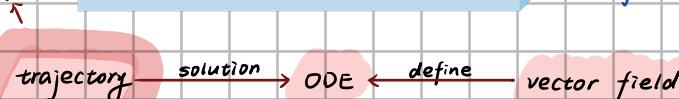
$$\psi: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d \quad (x_0, t) \mapsto \psi_t(x_0)$$

$$\frac{d}{dt} \psi_t(x_0) = u(\psi_t(x_0))$$

$$\psi_0(x_0) = x_0$$

⇒ flow ODE

⇒ flow initial condition



Existence and Uniqueness Theorem ODEs

if the vector field $u_t(x)$ is continuously differentiable with bounded derivatives, then a unique solution to the ODE

$$x_0 = x_0, \quad \frac{dx}{dt} = u_t(x_t)$$

exist. In other words, a flow map exists.

Example: Linear ODE

$$u_t(x) = -\theta x \quad (\theta > 0)$$

$$\psi_t(x_0) = \exp(-\theta t) x_0$$

Proof:

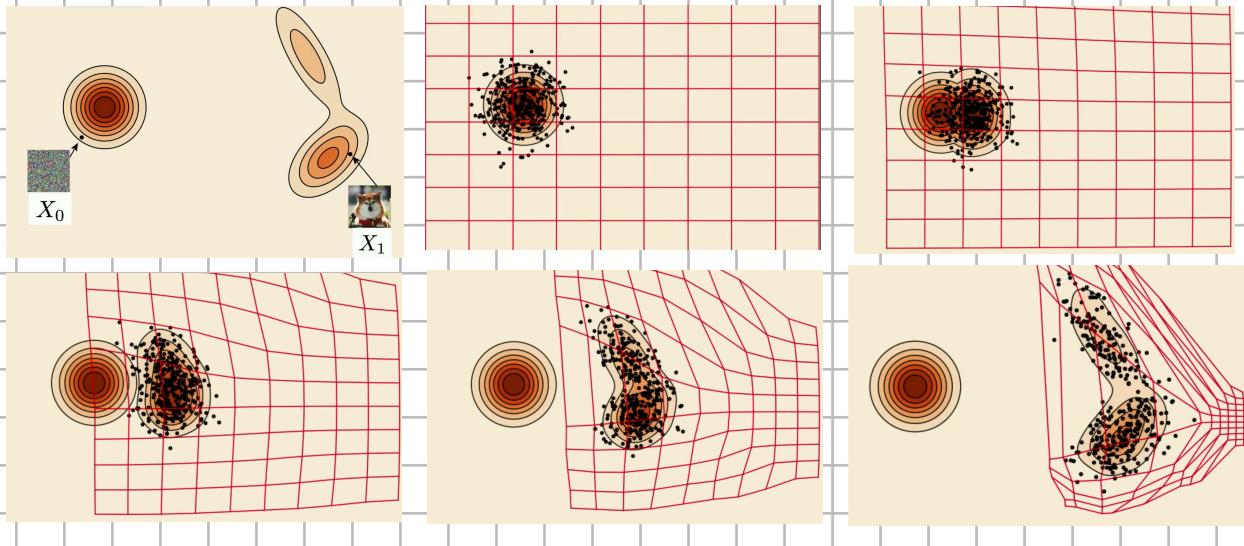
$$\begin{aligned} \text{initial condition: } \psi_t(x_0) &= \exp(-\theta t) x_0 = x_0 \\ \text{ODE: } \frac{d}{dt} \psi_t(x_0) &= \frac{d}{dt} [\exp(-\theta t) x_0] = -\theta \exp(-\theta t) x_0 \\ &= -\theta \psi_t(x_0) = u_t(\psi_t(x_0)) \end{aligned}$$

$$p_{\text{init}} \xrightarrow{\text{ODE}} p_{\text{data}}$$

Neural network: make the vector field a neural network

$$u_t^\theta: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d \quad \theta \text{ is parameters}$$

$$\begin{array}{l} \text{Random init: } x_0 \sim p_{\text{init}} \\ \text{ODE: } \frac{d}{dt} x_t = u_t^\theta(x_t) \end{array} \quad \text{goal: } x_1 \sim p_{\text{data}}$$



Algorithm 1 Sampling from a Flow Model with Euler method

Require: Neural network vector field u_t^θ , number of steps n

- 1: Set $t = 0$
- 2: Set step size $h = \frac{1}{n}$
- 3: Draw a sample $X_0 \sim p_{\text{init}}$
- 4: **for** $i = 1, \dots, n$ **do**
- 5: $X_{t+h} = X_t + h u_t^\theta(X_t)$
- 6: Update $t \leftarrow t + h$
- 7: **end for**
- 8: **return** X_1

Diffusion Models

SDEs (Stochastic differential equations)

ODE: deterministic trajectory \Rightarrow SDE: stochastic trajectory \Rightarrow stochastic process

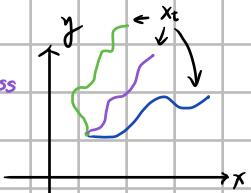
x_t : random variable ($t \in [0, 1]$)

$x: [0, 1] \rightarrow \mathbb{R}^d$, $t \rightarrow x_t$ is a random trajectory for every draw of x

fundamental object: ① vector field: $u: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$

② Diffusion coefficient: $\sigma: [0, 1] \rightarrow \mathbb{R}$ σ is a function of t : σ_t

injecting randomness or stochasticity into ODE



$x_0 = x_0$ (initial condition)

stochastic in SDE and fixed in ODE

$$dx_t = u_t(x_t) dt + \sigma_t dW_t$$

direction of a vector field inject noise

Brownian motion: $W = (W_t)_{t \geq 0}$

$$\textcircled{1} W_0 = 0 \quad W_t \in \mathbb{R}^d$$

$$\textcircled{2} \text{ Gaussian increments: } W_t - W_s \sim N(0, (t-s)I_d) \quad (0 \leq s < t)$$

$\textcircled{3} \text{ Independent increments: } W_{t_0} - W_{t_0}, \dots, W_{t_n} - W_{t_{n-1}}$ are independent
 $0 \leq t_0 \leq t_1 \leq \dots \leq t_n$

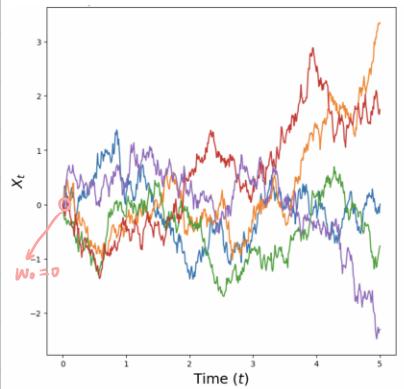
" dX_t " notation:

$$\text{ODE: } \frac{d}{dt} X_t = u_t(X_t) \Leftrightarrow X_{t+h} = X_t + h u_t(X_t) + h R_t(h) \rightarrow o(h), \text{ error term}$$

$$\lim_{h \rightarrow 0} R_t(h) = 0$$

$$\text{SDE: } dX_t = u_t(X_t) dt + \sigma_t dW_t \Leftrightarrow X_{t+h} = X_t + h u_t(X_t) + \sigma_t (W_{t+h} - W_t) + h R_t(h)$$

$$\lim_{h \rightarrow 0} \sqrt{\mathbb{E}[||R_t(h)||^2]} = 0 \quad \uparrow \quad \sim N(0, h I_d)$$



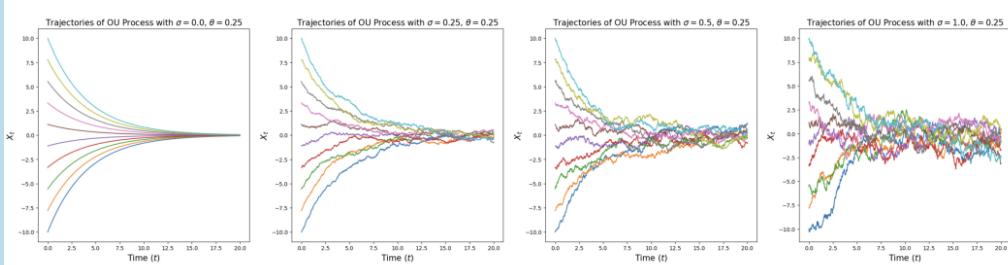
Existence and Uniqueness Theorem SDEs.

If the vector field $u_t(x)$ is continuously differentiable with bounded derivatives and the diffusion coeff. is continuous, then a unique solution to the SDE $X_0 = x_0, dX_t = u_t(X_t) dt + \sigma_t dW_t$ exists.

Example: Ornstein-Uhlenbeck Process:

a constant diffusion coefficient $\sigma_t = \sigma \geq 0$, a constant linear drift $u_t(x) = -\theta x$ ($\theta > 0$)

$$dX_t = -\theta X_t dt + \sigma dW_t$$



$$p_{\text{init}} \xrightarrow{\text{SDE}} p_{\text{data}}$$

Neural network $u_t^\theta: \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$

Diffusion coefficient: σ_t (Fixed)

random init: $X_0 \sim p_{\text{init}}$

goal: $x_t \sim p_{\text{data}}$

$$\text{SDE: } dX_t = u_t^\theta(X_t) dt + \sigma_t dW_t$$

Algorithm 2 Sampling from a Diffusion Model (Euler-Maruyama method)

Require: Neural network u_t^θ , number of steps n , diffusion coefficient σ_t

- 1: Set $t = 0$
 - 2: Set step size $h = \frac{1}{n}$
 - 3: Draw a sample $X_0 \sim p_{\text{init}}$
 - 4: **for** $i = 1, \dots, n$ **do**
 - 5: Draw a sample $\epsilon \sim \mathcal{N}(0, I_d)$
 - 6: $X_{t+h} = X_t + h u_t^\theta(X_t) + \sigma_t \sqrt{h} \epsilon$
 - 7: Update $t \leftarrow t + h$
 - 8: **end for**
 - 9: **return** X_1
-